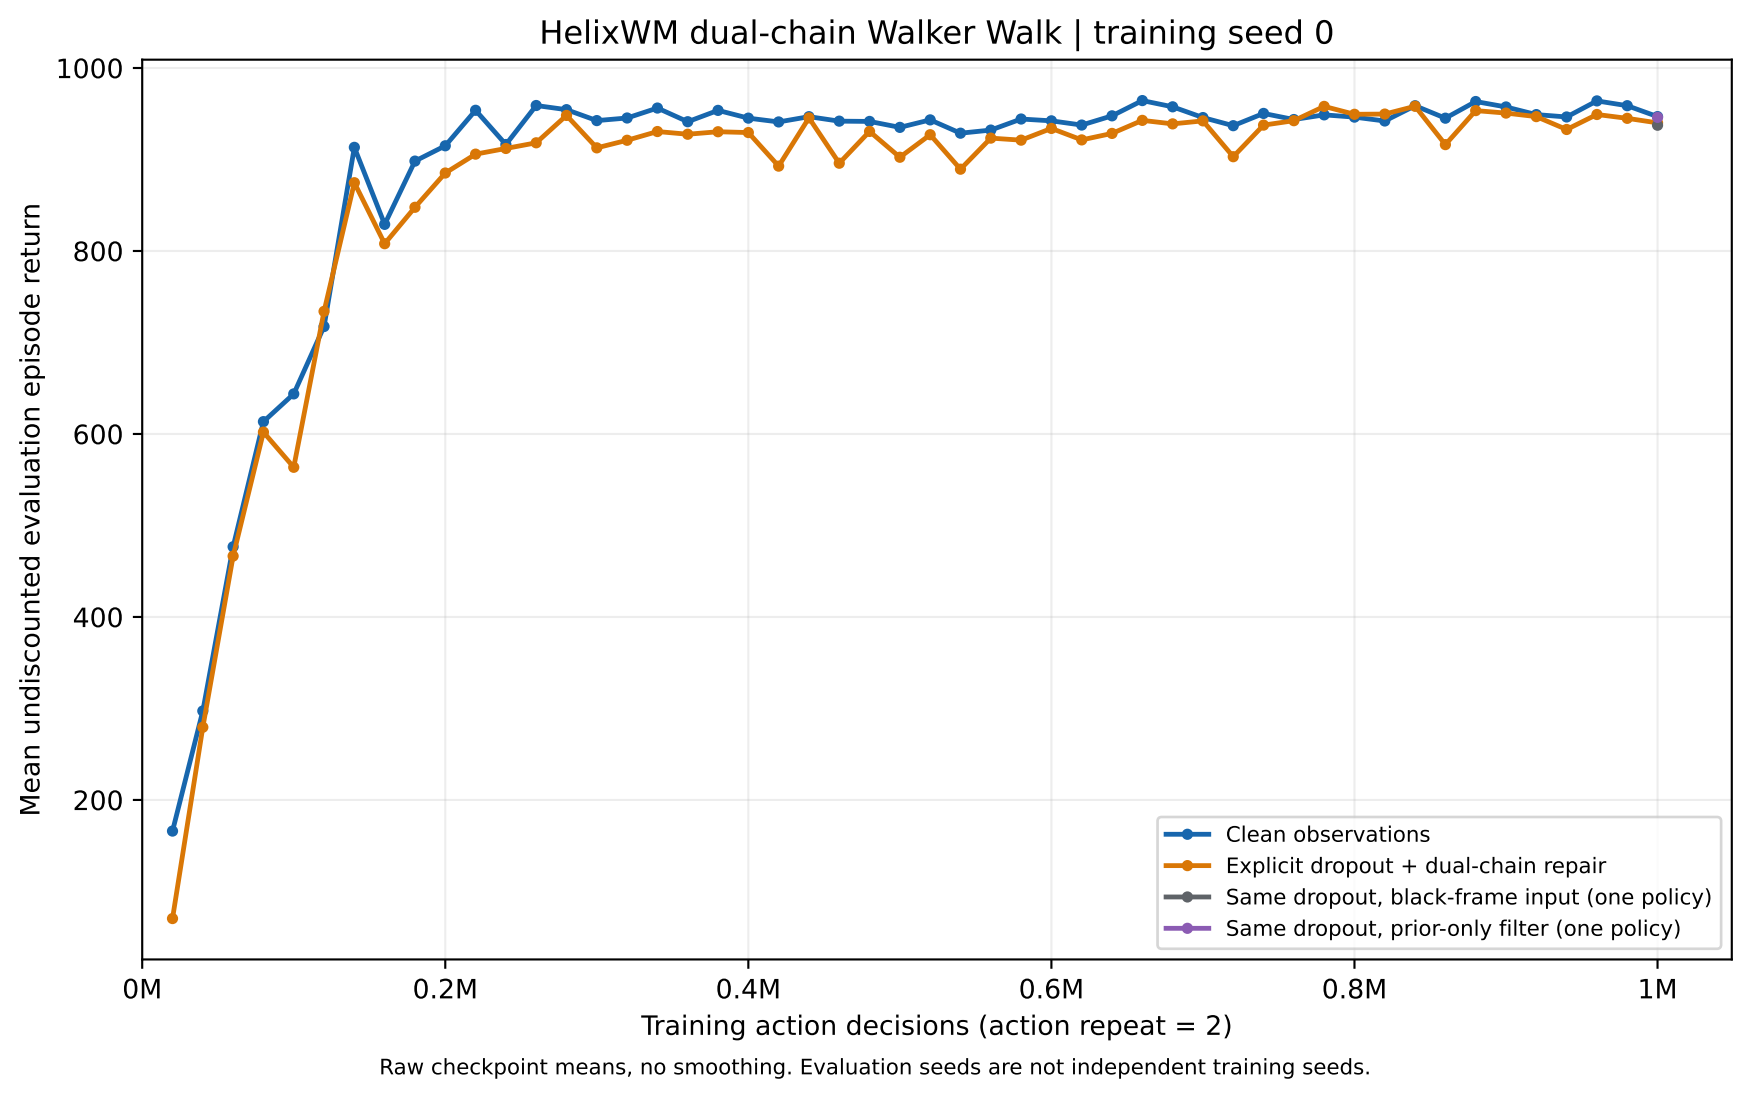
HelixWM dual-chain Walker Walk | training seed 0
Mean undiscounted evaluation episode return
1000
800
600
400
200
0M
0.2M
0.4M
0.6M
0.8M
1M
Training action decisions (action repeat = 2)
Clean observations
Explicit dropout + dual-chain repair
Same dropout, black-frame input (one policy)
Same dropout, prior-only filter (one policy)
Raw checkpoint means, no smoothing. Evaluation seeds are not independent training seeds.